

CS 224S: TensorFlow Tutorial

Lecture and Live Demo

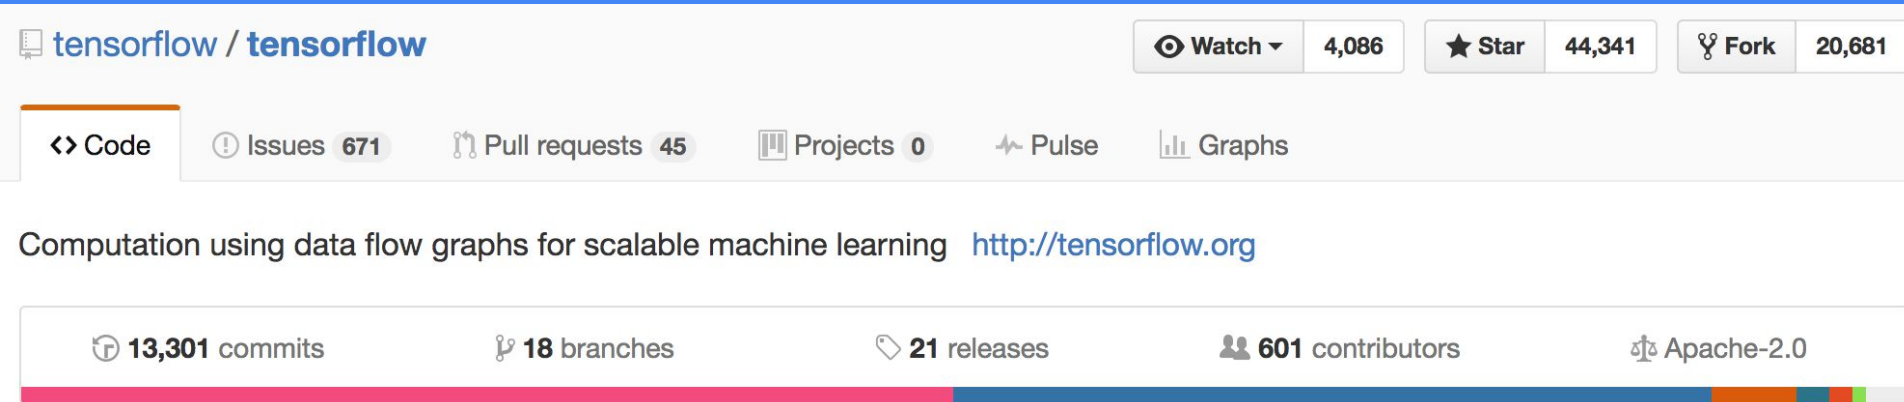
Pujun Bhatnagar
Rishabh Bhargava

Intro to Deep Learning Frameworks

- Scales machine learning code
- Computes gradients!
- Standardizes machine learning applications for sharing
- Zoo of Deep Learning frameworks available with different advantages, paradigms, levels of abstraction, programming languages, etc
- Interface with GPUs for parallel processing

In some ways, rightfully gives Deep Learning its name as a separate *practice*

What is TensorFlow?



The screenshot shows the GitHub repository for TensorFlow. At the top, the repository name 'tensorflow / tensorflow' is displayed. To the right, there are buttons for 'Watch' (4,086), 'Star' (44,341), and 'Fork' (20,681). Below these, there are tabs for 'Code', 'Issues' (671), 'Pull requests' (45), 'Projects' (0), 'Pulse', and 'Graphs'. The repository description is 'Computation using data flow graphs for scalable machine learning' with a link to 'http://tensorflow.org'. At the bottom, there are statistics: '13,301 commits', '18 branches', '21 releases', '601 contributors', and 'Apache-2.0' license. A progress bar is visible at the very bottom.

tensorflow / tensorflow

Watch 4,086 Star 44,341 Fork 20,681

<> Code Issues 671 Pull requests 45 Projects 0 Pulse Graphs

Computation using data flow graphs for scalable machine learning <http://tensorflow.org>

13,301 commits 18 branches 21 releases 601 contributors Apache-2.0

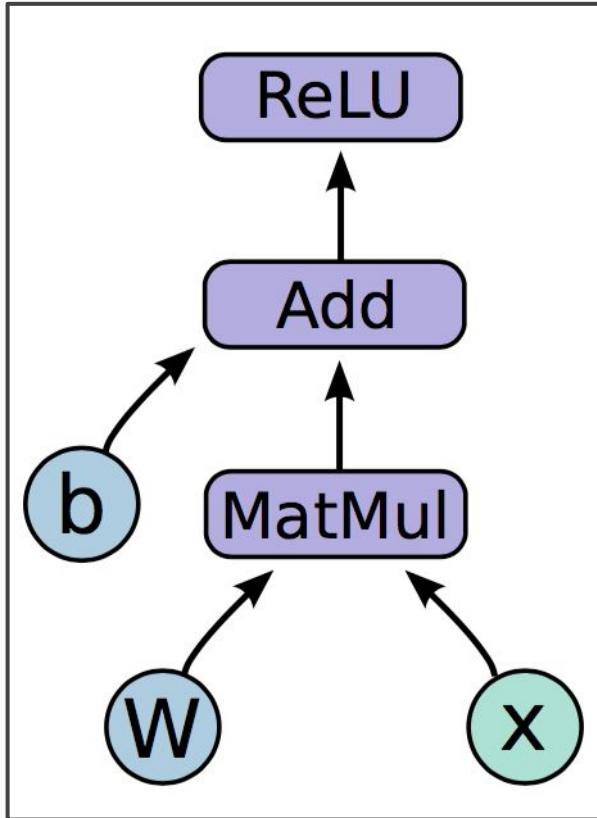
- Open source software library for numerical computation using data flow graphs
- Originally developed by Google Brain Team to conduct machine learning research
- “Tensorflow is an interface for expressing machine learning algorithms, and an implementation for executing such algorithms”

Programming model

Big idea: express a numeric computation as a **graph**.

- Graph nodes are **operations** which have any number of inputs and outputs
- Graph edges are **tensors** which flow between nodes

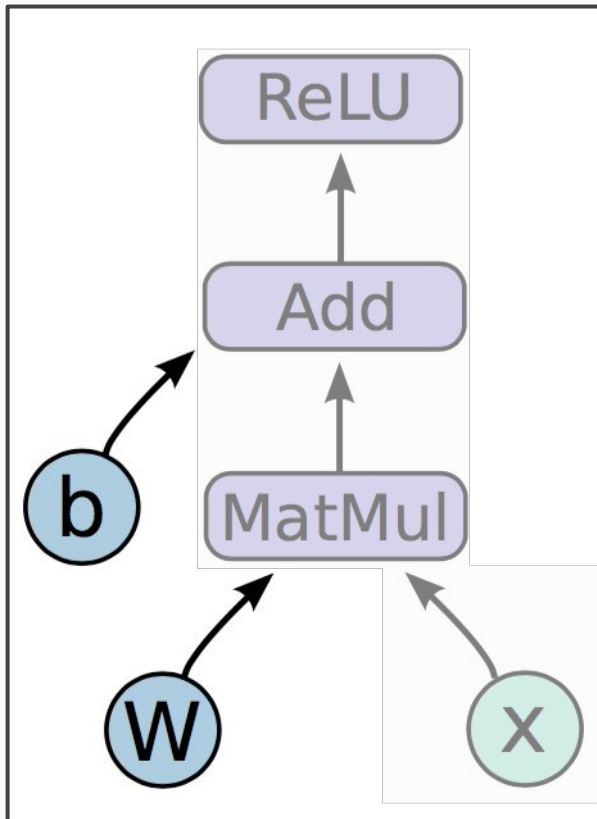
$$h = \text{ReLU}(Wx + b)$$



$$h = \text{ReLU}(Wx + b)$$

Variables are stateful nodes which output their current value.
State is retained across multiple executions of a graph

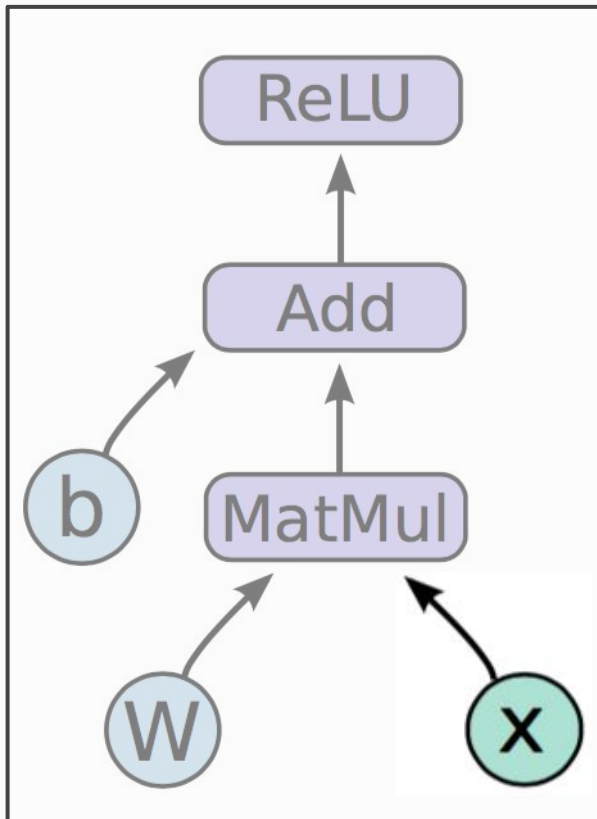
(mostly parameters)



$$h = \text{ReLU}(Wx + b)$$

Placeholders are nodes whose value is fed in at execution time

(inputs, labels, ...)



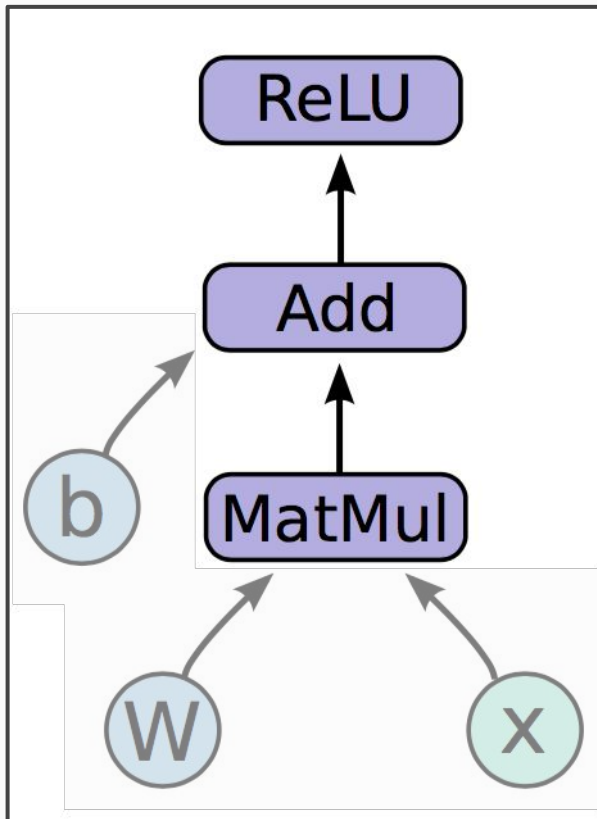
$$h = \text{ReLU}(Wx + b)$$

Mathematical operations:

MatMul: Multiply two matrix values.

Add: Add elementwise (with broadcasting).

ReLU: Activate with elementwise rectified linear function.



In code,

1. Create weights, including initialization

$$W \sim \text{Uniform}(-1, 1); b = 0$$

2. Create input placeholder x
 $m * 784$ input matrix

3. Build flow graph

```
import tensorflow as tf
```

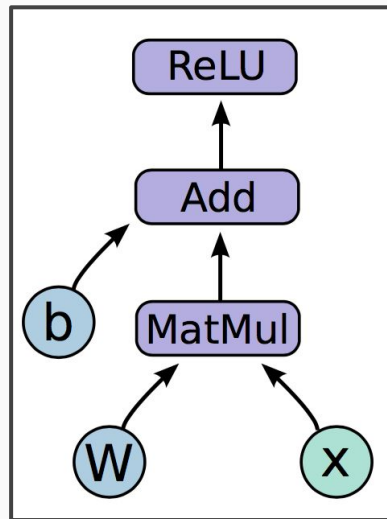
```
b = tf.Variable(tf.zeros((100,)))
```

```
W = tf.Variable(tf.random_uniform((784, 100), -1, 1))
```

```
x = tf.placeholder(tf.float32, (100, 784))
```

```
h = tf.nn.relu(tf.matmul(x, W) + b)
```

$$h = \text{ReLU}(Wx + b)$$



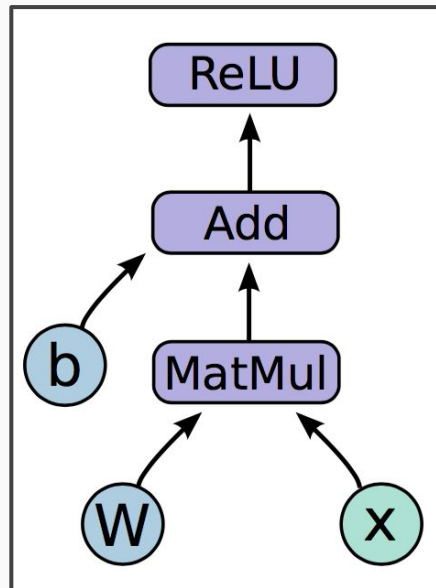
But where is the graph?

New nodes are automatically built into the underlying graph!
`tf.get_default_graph().get_operations():`

zeros/shape
zeros/Const
zeros
Variable
Variable/Assign
Variable/read
random_uniform/shape
random_uniform/min
random_uniform/max
random_uniform/RandomUniform

random_uniform/sub
random_uniform/mul
random_uniform
Variable_1
Variable_1/Assign
Variable_1/read
Placeholder
MatMul
add
Relu == h


h refers to an op!



How do we run it?

So far we have defined a **graph**.

We can deploy this graph with a **session**:
a binding to a particular execution
context (e.g. CPU, GPU)



Getting output

```
sess.run(fetches, feeds)
```

Fetches: List of graph nodes.

Return the outputs of these nodes.

Feeds: Dictionary mapping from graph nodes to concrete values. Specifies the value of each graph node given in the dictionary.

```
import numpy as np
import tensorflow as tf
```

```
b = tf.Variable(tf.zeros((100,)))
W = tf.Variable(tf.random_uniform((784, 100),
                                  -1, 1))
```

```
x = tf.placeholder(tf.float32, (100, 784))
h = tf.nn.relu(tf.matmul(x, W) + b)
```

```
sess = tf.Session()
sess.run(tf.initialize_all_variables())
sess.run(h, {x: np.random.random(100, 784)}))
```

So what have we covered so far?

We first built a **graph** using **variables** and **placeholders**

We then deployed the graph onto a **session**, which is the **execution environment**

Next we will see how to **train** the **model**

How do we define the loss?

Use **placeholder** for **labels**

Build loss node using labels and **prediction**

```
prediction = tf.nn.softmax(...) #Output of neural network
label = tf.placeholder(tf.float32, [100, 10])

cross_entropy = -tf.reduce_sum(label * tf.log(prediction), axis=1)
```

How do we compute Gradients?

```
train_step = tf.train.GradientDescentOptimizer(0.5).minimize(cross_entropy)
```

- `tf.train.GradientDescentOptimizer` is an **Optimizer** object
- `tf.train.GradientDescentOptimizer(lr).minimize(cross_entropy)` adds optimization **operation** to computation graph

TensorFlow graph **nodes** have **attached gradient operations**

Gradient with respect to **parameters** computed with **backpropagation**

...automatically

Creating the train_step op

```
prediction = tf.nn.softmax(...)
label = tf.placeholder(tf.float32, [None, 10])

cross_entropy = tf.reduce_mean(-tf.reduce_sum(label * tf.log(prediction),
reduction_indices=[1]))

train_step = tf.train.GradientDescentOptimizer(0.5).minimize(cross_entropy)
```


Variable sharing: naive way

```
variables_dict = {  
    "weights": tf.Variable(tf.random_normal([784, 100]),  
                           name="weights"),  
    "biases": tf.Variable(tf.zeros([100]), name="biases")  
}
```

Not good for encapsulation!

What's in a Name?

`tf.variable_scope()` provides simple name-spacing to avoid clashes


`tf.get_variable()` creates/accesses variables from within a variable scope

```
with tf.variable_scope("foo"):
    v = tf.get_variable("v", shape=[1]) # v.name == "foo/v:0"


with tf.variable_scope("foo", reuse=True):
    v1 = tf.get_variable("v") # Shared variable found!

with tf.variable_scope("foo", reuse=False):
    v1 = tf.get_variable("v") # CRASH foo/v:0 already exists!
```

Live Demo: Learning the MNIST dataset




RNNs



Cell types:

- Vanilla
- LSTM
- GRU

RNN sequences



RNN + Tensorflow

- Most calculations are performed on a per-batch basis
- RNNs expects tensors of shape $[B, T, \dots]$ as input, where B is batch size, T is length in time of each input.
- **What if each sequence is not the same length T ?**
- Solution: **Batching and Padding**

Batching and Padding

- Imagine one of your sequences is of length 1000, but the average sequence length is 20.
 - Pad all sequences to length 1000 – huge waste of memory and time!
 - Make batches of size 32, and pad all examples in the batch to maximum sequence length in batch. Only one batch suffers!
 - Padding: Appending 0's to shorter sequences in a batch to make them equal length.
- Dynamically unroll the RNN (example on next slide)

dynamic_rnn

```
# Create input data
X = np.random.randn(2, 10, 8)

# The second example is of length 6
X[1,6:] = 0
X_lengths = [10, 6]

cell = tf.nn.rnn_cell.LSTMCell(num_units=64, state_is_tuple=True)

outputs, last_states = tf.nn.dynamic_rnn(
    cell=cell,
    dtype=tf.float64,
    sequence_length=X_lengths,
    inputs=X)
```

In Summary:

1. Build a graph
 - a. Feedforward / Prediction
 - b. Optimization (gradients and train_step operation)
2. Initialize a session
3. Train with `session.run(train_step, feed_dict)`

Questions?

Acknowledgments

Barak Oshri, Nishith Khandwala, CS224N CAs last quarter

Jon Gauthier, Natural Language Processing Group, Symbolic Systems

Bharath Ramsundar, PhD Student, Drug Discovery Research

Chip Huyen, Undergraduate, teaching CS20SI: TensorFlow for Deep Learning Research last quarter

Live Demo: Tackling MNIST using Tensorflow

Insert link to the ipython notebook

Link to the github repo: <https://github.com/pbhatnagar3/cs224s-tensorflow-tutorial>